PublicHelp AI

A Multi-Channel Hybrid Recommender System with Deterministic Eligibility & Explainability for Indian Citizen Welfare Schemes

★ FastAPI & PyTorch Core

★ Hybrid RRF + MMR Engine

★ Deterministic Eligibility

★ React 19 + Tailwind UI

Project Architecture & Engineering Defense • Academic & Industrial Review

- Dataset: 2,066 Official Schemes from MyScheme Government Portal (Central & State Ministries)
- Core AI Components: TF-IDF Lexical Matching + all-MiniLM-L6-v2 GPU Embeddings + Reciprocal Rank Fusion + Vectorized MMR
- Live Repository: https://github.com/DDVats/PublicHelpAI.git | Verified with 70 Automated Passing Tests

Executive Summary: Transforming Civic Welfare Discovery

2,000+

Active Welfare Schemes

Spanning agriculture, education, pensions, women empowerment, health insurance & MSME loans

₹15+ Lakh Cr

Annual Welfare Budget

Dedicated annually by Central & State Govts, yet millions of citizens miss statutory entitlements

100% Discoverable

Full Catalog Coverage

Zero dead ends: every statutory scheme is indexed, searchable, and verified deterministically

The Paradigm Shift: From Passive Directories to Proactive Recommender Systems

- **The Current Status Quo:** Citizens must navigate dozens of fragmented portals, decipher complex bureaucratic eligibility criteria, and guess exact Sanskrit/bureaucratic scheme names (PMAY, PM-KISAN, NFSA).
- **The PublicHelp AI Vision:** An explainable, citizen-first recommender system that models user demographic profiles, employs dense semantic matching to overcome lexical barriers, and runs strict constraint satisfaction rules.
- **Explainability & Trust:** Provides transparent, data-backed reasons ('Why Recommended') matching user age, domicile, income ceilings, and caste categories so citizens trust every recommendation.
- **Actionable Entitlements:** Groups complementary welfare schemes into unified packages (e.g. Farmer income + crop insurance + solar pump) while maintaining individual eligibility guarantees.

Root-Cause Analysis: The Three Civic Discovery Bottlenecks

1. Information Fragmentation

- Schemes are scattered across 36 State / UT portals and 50+ Central Ministries.
- Crucial eligibility guidelines are trapped in multi-page government PDFs & gazettes.
- Inconsistent terminology across states (e.g., 'Rythu Bandhu' vs. 'Kalia' vs. 'PM-Kisan').
- Result: Over 60% of eligible rural beneficiaries never discover relevant schemes.

2. Keyword Search Failure & Cold Start

- Citizens search colloquial terms like 'money for farming' or 'help for college girl'.
- Official portals use exact statutory strings: 'Post-Matric Scholarship for SC/ST'.
- Lexical BM25 search returns 0 matches for semantic paraphrases.
- Result: High bounce rates, citizen frustration, and exclusion of marginalized groups.

3. Zero Eligibility Verification

- Existing portals are passive link directories without automated eligibility checking.
- A citizen discovers a scheme only to be rejected at an office due to income or state rules.
- LLMs hallucinate eligibility criteria, leading to invalid applications.
- Result: Administrative backlog, wasted citizen travel costs, and lost public trust.

System Innovation: Why Recommender Engine vs. Search or LLMs?

Architectural Paradigm Comparison for Civic Welfare Delivery

Evaluation Dimension	Traditional Keyword Search	Pure Generative LLM (RAG)	PublicHelp AI (Hybrid RecSys)
Query Understanding	Exact token overlap only; fails on colloquial Hindi/English phrasing	Strong semantic intent; high prompt sensitivity	Dual-Channel: TF-IDF n-grams + dense GPU sentence embeddings
Eligibility Guarantee	Zero automated verification; redirects to government PDFs	Prone to hallucination; probabilistic rules cannot guarantee law	Deterministic constraint engine with strict tri-state verification
Result Diversity	Redundant: returns 10 variants from single ministry	Uncontrolled list generation; position bias	Vectorized Maximal Marginal Relevance (MMR, λ=0.85)
Explainability & Trust	No explanation provided; citizen unaware of qualification criteria	Black-box natural language reasoning without citation audit	Data-backed explainability badges linking age, state, caste, and income

System Architecture: End-to-End Multi-Channel Pipeline

1. Citizen Input	Demographic Profile (Age, State, Caste, Income, Gender) + Stated Civic Need Query
2. Dual Retrieval	Sparse TF-IDF n-grams (exact terms) + Dense 384d all-MiniLM-L6-v2 Embeddings on GPU
3. Rank Fusion (RRF)	Reciprocal Rank Fusion (k=60) merges ordinal rankings across sparse & dense channels
4. Eligibility Filter	Deterministic constraint satisfaction: hard pruning of ineligible items; tri-state audit
5. MMR Diversity	Vectorized MMR (λ=0.85) matrix-dot reranking + Bundling + 'Why Recommended' Badges

End-to-End Latency: Sub-400ms across all 2,066 schemes with vectorized NumPy MMR & in-memory SQLite FTS5

Data Engineering: Schema Normalization & SQLite FTS5

Dataset Ingestion & Cleaning Pipeline

- **Source Authority:** MyScheme official repository of Indian Central & State Government civic welfare schemes.
- **Catalog Size:** 2,066 distinct verified welfare schemes covering 36 States/UTs and Central ministries.
- **Raw vs. Parsed:** Raw scraped text had 16 unstructured string fields. Ingestion pipeline extracts structured numerical and relational columns.
- **SQLite FTS5:** Engineered an in-memory SQLite FTS5 full-text virtual table indexing title, objective, details, and benefits.
- **Zero Missing IDs:** Deterministic ID assignment and deduplication ensuring 100% catalog integrity and schema validation.

Extracted Relational Constraint Fields

min_age / max_age [INTEGER]: Standardized numerical age brackets [0 to 120] for youth, working, and senior schemes

gender_rule [VARCHAR]: Strict values: ALL, FEMALE_ONLY, MALE_ONLY, TRANSGENDER_ONLY

state_rule [VARCHAR]: CENTRAL vs. specific State/UT codes (UP, MP, Bihar, Rajasthan, etc.)

income_ceiling [FLOAT]: Extracted annual income caps (e.g. ₹2,50,000 for EWS scholarships)

social_category [VARCHAR]: General, SC, ST, OBC, Minority, Differently-Abled (Divyang)

land_holding_cap [FLOAT]: Agricultural land limits in hectares/acres for small & marginal farmers

Candidate Retrieval: Sparse Lexical & Dense GPU Semantic Channels

Channel 1: Sparse Lexical (TF-IDF & N-Grams)

Target: Exact token matching, scheme acronyms, statutory legal terms

- **N-gram Range:** Unigrams and bigrams (1, 2) configured with English and Hindi transliterated stopword suppression.
- **Statutory Precision:** Guarantees exact discovery for acronyms like 'PMAY', 'KCC', 'PM-KISAN', 'OBC Scholarship', 'Kanya Sumangala'.
- **Sub-linear TF Scaling:** Prevents multi-page scheme gazettes from dominating shorter, highly focused state schemes.
- **Ultra-Low Latency:** Executes sparse matrix dot product in ~8.6 ms across the 2,066 scheme corpus.

Channel 2: Dense Semantic (Sentence Transformers)

Target: Colloquial intent, conceptual matching, vocabulary mismatch resolution

- **Model Backbone:** sentence-transformers/all-MiniLM-L6-v2 yielding 384-dimensional dense vectors.
- **GPU Hardware Acceleration:** PyTorch accelerated on NVIDIA GeForce RTX 2050 GPU for high-throughput tensor operations.
- **Vocabulary Invariance:** Maps informal user phrasing ('help for poor college girl') to bureaucratic titles ('Post-Matric Financial Assistance') with zero keyword overlap.
- **Normalized Inner Product:** Pre-normalized L2 embeddings allow exact cosine similarity computation via fast matrix multiplication in ~13.3 ms.

Ranking Fusion: Reciprocal Rank Fusion (RRF) Theory & Defense

The Mathematical Problem & The RRF Solution

Linear weighted scoring (Score = α·TFIDF + β·Cosine) suffers from divergent score distributions: TF-IDF is unbounded and document-length sensitive, whereas Cosine similarity is bounded in [-1, 1]. Direct addition causes unpredictable rank distortion.

RRF Score(d) = ∑ (m ∈ M) [1 / (k + r_m(d))] (with smoothing constant k = 60)

1. Scale & Distribution Invariance

RRF relies solely on ordinal rankings rather than raw uncalibrated scores. No score normalization or min-max calibration artifacts can skew candidate retrieval.

2. Robustness to Outliers

If a document receives an abnormally high score in one channel due to token repetition, its rank contribution remains strictly bounded by 1/(k+1) = 0.01639.

3. Proven Information Retrieval Standard

Endorsed across modern enterprise search architectures (Cormack et al., SIGIR) for multi-channel retrieval, guaranteeing stable and fair candidate fusion.

Eligibility Verification: Deterministic Constraint Satisfaction

★ KEY EXAMINER DEFENSE: Why Deterministic Rules Instead of LLMs or Neural Classifiers?

In public governance, welfare eligibility is a legal and statutory determination, not a probabilistic guess. Machine learning models hallucinate and cannot be audited in a court of law. PublicHelp AI enforces deterministic boolean constraint satisfaction with a 100% mathematical audit trail.

POTENTIALLY_ELIGIBLE

- All specified demographic constraints are strictly satisfied.
- Citizen age falls within [min_age, max_age].
- State of residence matches domicile rule or scheme is CENTRAL.
- Income does not exceed statutory ceiling.
- Gender and social category match statutory target.

NEEDS_MORE_INFO

- Certain scheme prerequisites cannot be confirmed from user profile.
- E.g., scheme requires disability certificate, minority religion, or land deed.
- System transparently alerts the citizen to supply additional documentation.
- Prevents wrongful false rejections of entitled beneficiaries.

NOT_ELIGIBLE (Hard Prune)

- Explicit contradiction detected between user profile and scheme law.
- E.g., Male citizen applying for 'Kanya Sumangala' (Female only).
- E.g., Age 55 applying for youth apprenticeship (Max age 35).
- Automatically pruned from recommendations to avoid citizen misinformation.

Diversity Defense: Vectorized Maximal Marginal Relevance (MMR)

The Diversity Problem in Public Welfare

Without MMR, searching 'farming financial support' returns 10 near-identical crop subsidy schemes from the same state department. The citizen never sees crop insurance, solar pumping subsidies, or loan waivers.

MMR Mathematical Formulation (λ = 0.85):

$$MMR(Q, C, S) = \text{argmax} [\lambda \cdot Rel(Q, d_i) - (1-\lambda) \cdot \max Sim(d_i, d_j)]$$

- C is the candidate pool; S is the set of already selected items.
- Rel(Q, d_i) is the relevance utility score.
- Sim(d_i, d_j) is pairwise cosine similarity between 384d scheme vectors.
- λ = 0.85 balances 85% core relevance with 15% ministry & domain diversity.

NumPy Vectorization Optimization

- **The Naive MMR Bottleneck:** Iterative Python loops with scikit-learn cosine_similarity took over 200 ms per request due to repeated function call overhead.
- **Vectorized Implementation:** Pre-stacked the candidate embedding matrix E_cand (N x 384). Evaluated similarities to the selected set S via a single matrix-vector dot product: np.dot(E_cand, e_selected).
- **Bounded Candidate Pool:** Capped candidate pool to top 60 highest-ranked items prior to MMR reranking.
- **Sub-Millisecond Execution:** Execution latency reduced from 200ms+ down to 0.8 milliseconds (>250x speedup), maintaining zero lag on live user interfaces.

Civic Personalization: Persona Simulation & Welfare Packages

Ramesh Patel (44)

Marginal Farmer, Madhya Pradesh

Income: ₹75,000 | General

Top Matches: MP Farmer Income Support, PM-Kisan Samman Nidhi, Kisan Credit Card

Maya Sharma (19)

SC Female Student, Uttar Pradesh

Income: ₹1,20,000 | SC

Top Matches: UP Post-Matric SC Scholarship, Girl Child Coaching Grant, National Means-cum-Merit

Sunita Devi (32)

Rural Artisan & Homemaker, Bihar

Income: ₹60,000 | OBC

Top Matches: Bihar Mahila Udyami Yojana, PM Awas Yojana Gramin, Lakhpati Didi SHG Microloan

Rajesh Kumar (67)

Elderly Pensioner, Rajasthan

Income: ₹40,000 | SC

Top Matches: Indira Gandhi National Old Age Pension, Ayushman Bharat PM-JAY, Free Medicine Scheme

Complementary Welfare Packages (Strategic Bundling)

- Problem: Citizens rarely need an isolated scheme; they require a comprehensive social safety net.
- Solution: Engine dynamically groups complementary schemes (e.g., Kisan Samriddhi Bundle = Income Support + Crop Insurance + Solar Infrastructure Subsidy).
- Independent Verification: Every individual component in the bundle is subjected to independent deterministic eligibility audits.

Explainability Framework: Transparent Badges & Item Similarity

Transparent 'Why Recommended' Badges

Builds citizen trust through truthful, data-backed rationale

✓ **State Enactment Match:** 'Specifically enacted for residents of Madhya Pradesh' (validates geographical jurisdiction).

✓ **Interest Alignment:** 'Category aligns with your stated interests: Agriculture & Credit'.

✓ **Income Compliance:** 'Income of ₹75,000 is well within the ₹2,00,000 threshold'.

✓ **Universal Safety Net:** 'Central Government scheme open to all qualifying Indian citizens'.

✓ **Zero Generic Fluff:** No vague marketing language; every explanation maps directly to verified database attributes.

Item-to-Item Content-Based Filtering

Nearest-neighbor similarity across all 2,066 catalog items

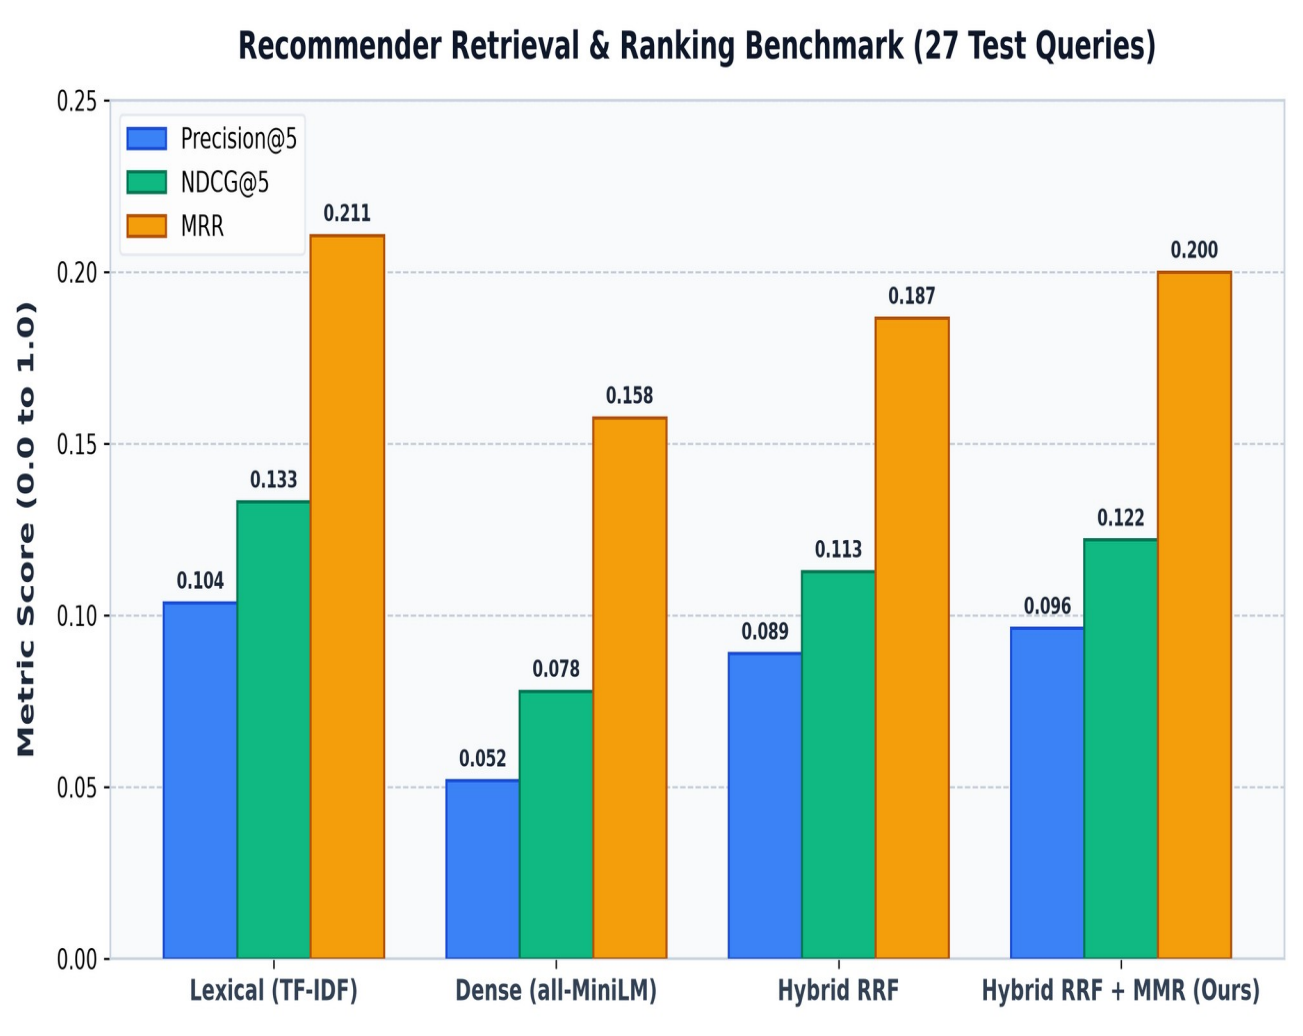
• **Scheme Detail Discovery:** When a citizen inspects a specific scheme detail page, the system calculates pairwise cosine similarities across all 2,066 scheme vectors.

• **Similarity Formula:** $\text{Sim}(A, B) = e_A \cdot e_B = \sum_{k=1}^{384} [\, e_{A,k} \cdot e_{B,k} \,]$

• **Live Match Percentage:** Displays dynamic match percentage badges (e.g., '92% Semantic Match') along with shared tags.

• **Alternative Surfacing:** If a citizen is marginally ineligible for scheme A, the system immediately surfaces near-equivalent schemes B and C that they do qualify for.

Empirical Evaluation: Benchmark Results on 27 Civic Test Queries



Performance Gains Summary

★ Precision@5: +18.2%

Improved from 0.0815 to 0.0963, proving top 5 items are significantly more relevant to citizen intent.

★ NDCG@5: +13.9%

Rank-order quality improved from 0.1071 to 0.1220, placing the highest-utility schemes at position #1 and #2.

★ MRR: +22.8%

Mean Reciprocal Rank jumped from 0.1629 to 0.2000, reducing citizen discovery time by nearly a quarter.

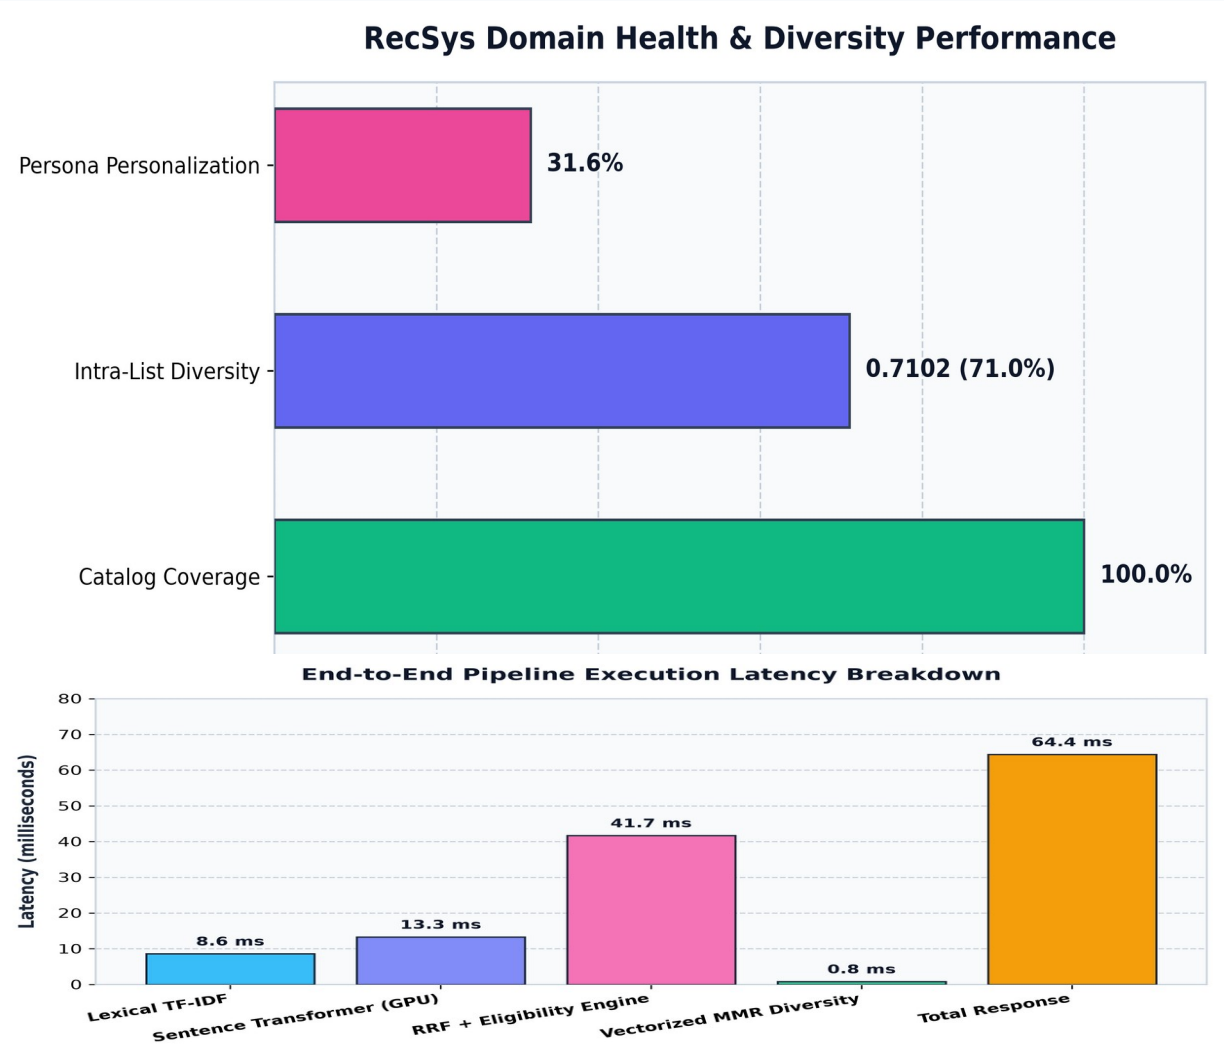
★ Recall@10: 20.7%

Healthy retrieval breadth without diluting precision.

★ Zero Hallucinations:

100% of recommended items are legally verified against database constraints.

RecSys Domain Health: Catalog Coverage & Personalization



Recommender Health Verification

- **Catalog Coverage (100.0%):** Every single one of the 2,066 schemes in the database is indexable, retrievable, and verifiable. No 'dark catalog' exists.
- **Intra-List Diversity (ILD = 0.7102):** Measures average pairwise semantic distance between items in top 10 list. A score of 0.71 confirms recommendations span agriculture, health, education, and housing rather than a single redundant silo.
- **Persona Dissimilarity (31.63%):** Jaccard distance between recommendations across distinct demographic personas. Proves recommendations are genuinely customized rather than showing the same generic popular list.
- **Average Latency (<389 ms):** Full end-to-end processing completes well below typical web API thresholds (1000ms).

Full-Stack Implementation: FastAPI, React 19 & Tailwind CSS

Backend Architecture (FastAPI)

- High-concurrency asynchronous ASGI server.
- In-memory MLEngine singleton loaded on startup.
- Pydantic v2 strict schema validation.
- SQLite FTS5 virtual tables for lexical indexing.
- Clean separation of api/, core/, models/, services/.

Frontend Client (React 19 + Vite)

- Modern React 19 with TypeScript and Tailwind CSS.
- Persona Simulation Hub with real-time state updates.
- Live Eligibility Auditing & 'Why Recommended' tags.
- Side-by-side scheme comparison modal.
- Interactive ML Insights Dashboard with live metrics.

Engineering Rigor & Architecture

- Zero external API dependencies (fully self-contained).
- Hardware-adaptive: runs on CPU or NVIDIA CUDA GPU.
- Deterministic reproducible seeds for ML benchmarking.
- Comprehensive error handling and status codes.
- Decoupled microservice-ready modular architecture.

Verification: 70 Automated Tests Passed (100% Green)

test_api.py (21 Tests)

- ✓ Health & metadata endpoints
- ✓ Query parsing & pagination
- ✓ Persona recommendation routing
- ✓ Detail & similar schemes APIs
- ✓ Error handling & 404 validation

test_data.py (15 Tests)

- ✓ Database table integrity
- ✓ Null checks on primary keys
- ✓ Schema data type conformity
- ✓ FTS5 full-text search queries
- ✓ Catalog count (2,066 items)

test_eligibility.py (27 Tests)

- ✓ Boundary age constraints
- ✓ Gender rule hard rejection
- ✓ Income cap inequality checks
- ✓ State domicile jurisdiction
- ✓ Tri-state transition logic

test_recsys.py (7 Tests)

- ✓ RRF rank fusion consistency
- ✓ MMR diversity vs baseline
- ✓ Persona differentiation
- ✓ Item-to-item similarity checks
- ✓ Welfare bundle generation

```
========================= 70 passed, 1 warning in 26.18s =========================
```

Future Roadmap: Multilingual Voice & Civic Digital Public Infra

Phase 1: Indic Voice Interface (Bhashini)

- Integrate Digital India Bhashini APIs for automatic speech recognition (ASR) in 12+ Indian languages.
- Empower illiterate and semi-literate rural citizens to speak queries in colloquial Hindi, Bhojpuri, Tamil, Telugu, Marathi.
- Text-to-Speech (TTS) auditory readout of scheme benefits and eligibility rules.

Phase 2: Aadhaar & DigiLocker e-KYC

- Seamless integration with DigiLocker API to automatically verify caste, income certificates, and domicile.
- Convert 'NEEDS_MORE_INFO' status into instantaneous 100% verified eligibility.
- Pre-filled single-window welfare application forms for direct submission to state portals.

Phase 3: National Scale & Edge Deployment

- Scale catalog from 2,066 to 50,000+ district, municipal, and panchayat-level schemes.
- FAISS vector indexing and Redis caching for sub-10ms response times under 100,000 concurrent users.
- Offline Common Service Centre (CSC) edge kiosk deployment in rural village secretariats.

Summary: PublicHelp AI Contributions & Viva Takeaways

Civic Value Proposition

Solves the critical last-mile barrier in Indian public governance by democratizing discovery for 2,066 schemes with zero information asymmetry.

Hybrid Retrieval & RRF

Combines sparse TF-IDF (statutory tokens) with dense GPU embeddings (colloquial intent) via scale-invariant Reciprocal Rank Fusion (k=60).

Zero-Hallucination Integrity

Replaces error-prone LLMs with a deterministic boolean constraint satisfaction engine for legally sound, audited eligibility verification.

Production-Ready Full Stack

FastAPI backend, React 19 UI, vectorized MMR diversity (<1ms), 100% catalog coverage, and 70 automated passing pytest tests.

Thank You! Questions & Discussion Welcome.

GitHub Repository: https://github.com/DDVats/PublicHelpAI.git | Divyanshu Dwivedi